

# (Shared)MLP\_Week\_4\_SWI(01\_10\_22)

October 28, 2022

You are working as a data scientist in a big automobile company. Your company aspires to enter the US market by setting up their manufacturing unit there and producing cars locally to give competition to their US and European counterparts.

They got some data to understand the factors on which the pricing of cars depends in the American market, since those may vary different from the indian market. The company wants to know:

Which variables are significant in predicting the price of a car, How well those variables describe the price of a car Based on various market surveys.

## 1 Business Goal:

Data science team are required to model the price of cars with the available independent variables. It will be used by the management to understand how exactly the prices vary with the independent variables. They can accordingly manipulate the design of the cars, the business strategy etc. to meet certain price levels. Further, the model will be a good way for management to understand the pricing dynamics of a new market.

## 2 Step-1: Importing Libraries

```
[1]: # Importing the libraries
import numpy as np
import pandas as pd
from numpy import math

from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score

import seaborn as sns
import matplotlib.pyplot as plt
import warnings
```

```
warnings.filterwarnings('ignore')
```

### 2.0.1 DataSet Information:

Dataset\_link: <https://drive.google.com/file/d/1EXZXCvI-be9M2zB7H1NkfDkSUSSt8poCc/view?usp=sharing>

#### **Features:-**

Car\_ID: Unique id of each observation (Integer)

Symboling: Its assigned insurance risk rating, A value of +3 indicates that the auto is risky, -3 that it is probably pretty safe.(Categorical)

CarName: Name of car company (Categorical)

fueltype: Car fuel type i.e gas or diesel (Categorical)

aspiration: Aspiration used in a car (Categorical)

doornumber: Number of doors in a car (Categorical)

carbody: body of car (Categorical)

drivewheel: type of drive wheel (Categorical)

enginelocation: Location of car engine (Categorical)

wheelbase: Wheelbase of car (Numeric)

carlength: Length of car (Numeric)

carwidth: Width of car (Numeric)

carheight: height of car (Numeric)

curbweight: The weight of a car without occupants or baggage. (Numeric)

engine type: Type of engine. (Categorical)

cylindernumber: cylinder placed in the car (Categorical)

enginesize: Size of car (Numeric)

fuelsystem: Fuel system of car (Categorical)

bore ratio: Bore ratio of car (Numeric)

stroke: Stroke or volume inside the engine (Numeric)

compressionratio: compression ratio of car (Numeric)

horsepower: Horsepower (Numeric)

peakrpm: car peak rpm (Numeric)

citympg: Mileage in city (Numeric)

highwaympg: Mileage on highway (Numeric)

price(Dependent variable): Price of car (Numeric)

```
[ ]:
```

### 3 Step-2: Loading the data

```
[ ]: from google.colab import drive
drive.mount('/content/drive')
```

```
[ ]: #pd.read_csv("/content/drive/MyDrive/Dataset/Car_price in US market.csv")
```

```
[ ]: from google.colab import files
files.upload()
```

```
[5]: # Importing the dataset
dataset = pd.read_csv('Car_price in US market.csv')
```

### 4 Step-3: Data Inspection

#### 5 Question set-1:

- (i) No of data point
- (ii) No of features
- (iii) No of categorical features
- (iv) No of numerical features
- (v) No of NA values
- (vi) List of all features
- (vii) What about duplicate data ?

```
[6]: dataset.shape[1]
```

```
[6]: 26
```

```
[7]: dataset.head(5)
```

```
[7]:
```

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	\
0	1	3	alfa-romero giulia	gas	std	two	
1	2	3	alfa-romero stelvio	gas	std	two	
2	3	1	alfa-romero Quadrifoglio	gas	std	two	
3	4	2	audi 100 ls	gas	std	four	
4	5	2	audi 100ls	gas	std	four	

	carbody	drivewheel	enginelocation	wheelbase	...	enginesize	\
0	convertible	rwd	front	88.6	...	130	
1	convertible	rwd	front	88.6	...	130	
2	hatchback	rwd	front	94.5	...	152	
3	sedan	fwd	front	99.8	...	109	
4	sedan	4wd	front	99.4	...	136	

	fuelsystem	boreratio	stroke	compressionratio	horsepower	peakrpm	citympg	\
0	mpfi	3.47	2.68		9.0	111	5000	21
1	mpfi	3.47	2.68		9.0	111	5000	21
2	mpfi	2.68	3.47		9.0	154	5000	19
3	mpfi	3.19	3.40		10.0	102	5500	24
4	mpfi	3.19	3.40		8.0	115	5500	18

	highwaympg	price
0	27	13495.0
1	27	16500.0
2	26	16500.0
3	30	13950.0
4	22	17450.0

[5 rows x 26 columns]

```
[8]: dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   car_ID                205 non-null    int64
1   symboling              205 non-null    int64
2   CarName               205 non-null    object
3   fueltype              205 non-null    object
4   aspiration            205 non-null    object
5   doornumber            205 non-null    object
6   carbody               205 non-null    object
7   drivewheel            205 non-null    object
8   enginelocation        205 non-null    object
9   wheelbase             205 non-null    float64
10  carlength             205 non-null    float64
11  carwidth              205 non-null    float64
12  carheight             205 non-null    float64
13  curbweight            205 non-null    int64
14  enginetype            205 non-null    object
15  cylindernumber        205 non-null    object
16  enginesize            205 non-null    int64
```

```

17  fuelsystem          205 non-null    object
18  boreratio           205 non-null    float64
19  stroke              205 non-null    float64
20  compressionratio    205 non-null    float64
21  horsepower          205 non-null    int64
22  peakrpm             205 non-null    int64
23  citympg             205 non-null    int64
24  highwaympg          205 non-null    int64
25  price               205 non-null    float64
dtypes: float64(8), int64(8), object(10)
memory usage: 41.8+ KB

```

```
[9]: dataset.isnull().sum()
```

```

[9]: car_ID          0
    symboling        0
    CarName          0
    fueltype         0
    aspiration        0
    doornumber        0
    carbody          0
    drivewheel        0
    enginelocation    0
    wheelbase         0
    carlength         0
    carwidth          0
    carheight         0
    curbweight        0
    enginetype        0
    cylindernumber    0
    enginesize        0
    fuelsystem        0
    boreratio         0
    stroke            0
    compressionratio  0
    horsepower        0
    peakrpm           0
    citympg           0
    highwaympg        0
    price             0
    dtype: int64

```

```
[10]: dataset.describe()
```

```

[10]:
   count  car_ID  symboling  wheelbase  carlength  carwidth  carheight  \
mean    103.000000    0.834146    98.756585    174.049268    65.907805    53.724878

```

std	59.322565	1.245307	6.021776	12.337289	2.145204	2.443522
min	1.000000	-2.000000	86.600000	141.100000	60.300000	47.800000
25%	52.000000	0.000000	94.500000	166.300000	64.100000	52.000000
50%	103.000000	1.000000	97.000000	173.200000	65.500000	54.100000
75%	154.000000	2.000000	102.400000	183.100000	66.900000	55.500000
max	205.000000	3.000000	120.900000	208.100000	72.300000	59.800000

	curbweight	enginesize	boreratio	stroke	compressionratio	\
count	205.000000	205.000000	205.000000	205.000000	205.000000	
mean	2555.565854	126.907317	3.329756	3.255415	10.142537	
std	520.680204	41.642693	0.270844	0.313597	3.972040	
min	1488.000000	61.000000	2.540000	2.070000	7.000000	
25%	2145.000000	97.000000	3.150000	3.110000	8.600000	
50%	2414.000000	120.000000	3.310000	3.290000	9.000000	
75%	2935.000000	141.000000	3.580000	3.410000	9.400000	
max	4066.000000	326.000000	3.940000	4.170000	23.000000	

	horsepower	peakrpm	citympg	highwaympg	price
count	205.000000	205.000000	205.000000	205.000000	205.000000
mean	104.117073	5125.121951	25.219512	30.751220	13276.710571
std	39.544167	476.985643	6.542142	6.886443	7988.852332
min	48.000000	4150.000000	13.000000	16.000000	5118.000000
25%	70.000000	4800.000000	19.000000	25.000000	7788.000000
50%	95.000000	5200.000000	24.000000	30.000000	10295.000000
75%	116.000000	5500.000000	30.000000	34.000000	16503.000000
max	288.000000	6600.000000	49.000000	54.000000	45400.000000

```
[11]: features=(dataset.columns)
```

```
[12]: features
```

```
[12]: Index(['car_ID', 'symboling', 'CarName', 'fueltype', 'aspiration',
          'doornumber', 'carbody', 'drivewheel', 'enginelocation', 'wheelbase',
          'carlength', 'carwidth', 'carheight', 'curbweight', 'enginetype',
          'cylindernumber', 'enginesize', 'fuelsystem', 'boreratio', 'stroke',
          'compressionratio', 'horsepower', 'peakrpm', 'citympg', 'highwaympg',
          'price'],
          dtype='object')
```

```
[13]: len(dataset[dataset.duplicated()])
```

```
[13]: 0
```

## 6 Step-5: Exploratory data analysis

Question Set 2:

- 1) Give list of all numeric features
- 2) List of all categorical features
- 3) Type of distribution your dependent variable follow.
- 4) Plot different graph (histogram,box-plot,scatter plot e.t.c ) for all independent variable to get some insight.
- 5) some scaling is needed or not.
- (5) Comment about different categorical feature.
- (6) Create a function which converts string into numerical e.g- {"four": 4, "two": 2}

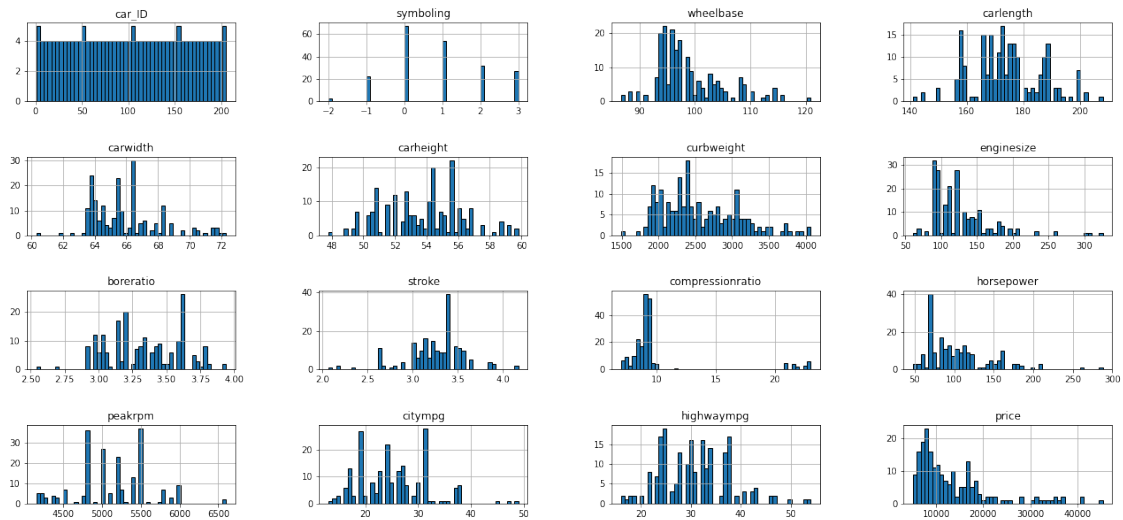
```
[14]: num_feat=dataset.describe().columns
      list(num_feat)
```

```
[14]: ['car_ID',
      'symboling',
      'wheelbase',
      'carlength',
      'carwidth',
      'carheight',
      'curbweight',
      'enginesize',
      'boreratio',
      'stroke',
      'compressionratio',
      'horsepower',
      'peakrpm',
      'citympg',
      'highwaympg',
      'price']
```

```
[15]: all_feat=list(dataset.columns)
      cat_feat= [i for i in all_feat if i not in num_feat]
      cat_feat
```

```
[15]: ['CarName',
      'fueltype',
      'aspiration',
      'doornumber',
      'carbody',
      'drivewheel',
      'enginelocation',
      'enginetype',
      'cylindernumber',
      'fuelsystem']
```

```
[ ]: dataset.hist(figsize=(22, 10), bins=50, edgecolor="black")
plt.subplots_adjust(hspace=0.7, wspace=0.4)
```




## 7 Distribution of independent variable

```
[ ]: plt.figure(figsize=(7,7))
sns.distplot(dataset['price'], color="g")
```


```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7f0cc17c6450>
```





```
[ ]: plt.figure(figsize=(7,7))  
sns.distplot(np.log10(dataset['price']),color="b")
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7f0cc0712ad0>
```



## 8 Correlation between different features

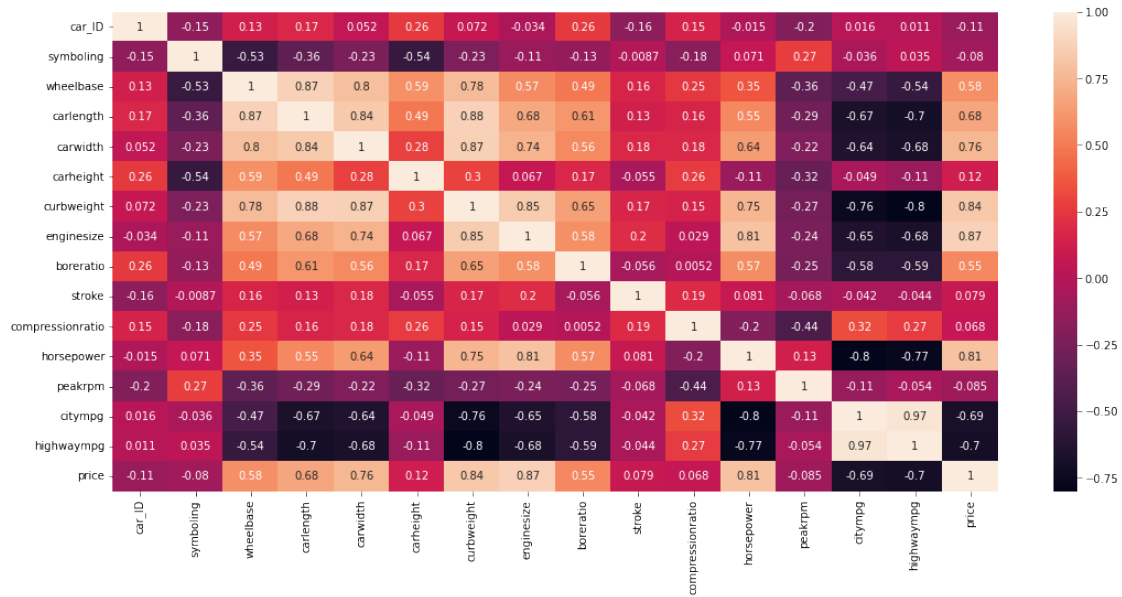
### Question Set 3:

- 1) Create a dataframe named `vif_df` which has 2 columns named variables and their vif.
- 2) Create a function named `make_vif_df` which returns the above dataframe.
- 3) Start removing column (whose vif > 5 or maybe 10), one at a time and check.
- 4) Apply one hot encoding on features like ("carbody", "enginetype", "fuelsystem")

```
[16]: plt.figure(figsize=(18,8))  
corre = dataset.corr()
```

```
sns.heatmap(corre,annot=True)
```

```
[16]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe5581eff90>
```



## 9 Feature engineering

- Check How many unique car brands are working in the USA?
- What is avg price of Jaguar car.

```
[ ]: dataset['CarName'].unique()
```

```
[ ]: array(['alfa-romero giulia', 'alfa-romero stelvio',
'alfa-romero Quadrifoglio', 'audi 100 ls', 'audi 100ls',
'audi fox', 'audi 5000', 'audi 4000', 'audi 5000s (diesel)',
'bmw 320i', 'bmw x1', 'bmw x3', 'bmw z4', 'bmw x4', 'bmw x5',
'chevrolet impala', 'chevrolet monte carlo', 'chevrolet vega 2300',
'dodge rampage', 'dodge challenger se', 'dodge d200',
'dodge monaco (sw)', 'dodge colt hardtop', 'dodge colt (sw)',
'dodge coronet custom', 'dodge dart custom',
'dodge coronet custom (sw)', 'honda civic', 'honda civic cvcc',
'honda accord cvcc', 'honda accord lx', 'honda civic 1500 gl',
'honda accord', 'honda civic 1300', 'honda prelude',
'honda civic (auto)', 'isuzu MU-X', 'isuzu D-Max ',
'isuzu D-Max V-Cross', 'jaguar xj', 'jaguar xf', 'jaguar xk',
'maxda rx3', 'maxda glc deluxe', 'mazda rx2 coupe', 'mazda rx-4',
```

```
'mazda glc deluxe', 'mazda 626', 'mazda glc', 'mazda rx-7 gs',
'mazda glc 4', 'mazda glc custom l', 'mazda glc custom',
'buick electra 225 custom', 'buick century luxus (sw)',
'buick century', 'buick skyhawk', 'buick opel isuzu deluxe',
'buick skylark', 'buick century special',
'buick regal sport coupe (turbo)', 'mercury cougar',
'mitsubishi mirage', 'mitsubishi lancer', 'mitsubishi outlander',
'mitsubishi g4', 'mitsubishi mirage g4', 'mitsubishi montero',
'mitsubishi pajero', 'Nissan versa', 'nissan gt-r', 'nissan rogue',
'nissan latio', 'nissan titan', 'nissan leaf', 'nissan juke',
'nissan note', 'nissan clipper', 'nissan nv200', 'nissan dayz',
'nissan fuga', 'nissan otti', 'nissan teana', 'nissan kicks',
'peugeot 504', 'peugeot 304', 'peugeot 504 (sw)', 'peugeot 604sl',
'peugeot 505s turbo diesel', 'plymouth fury iii',
'plymouth cricket', 'plymouth satellite custom (sw)',
'plymouth fury gran sedan', 'plymouth valiant', 'plymouth duster',
'porsche macan', 'porcshce panamera', 'porsche cayenne',
'porsche boxter', 'renault 12tl', 'renault 5 gtl', 'saab 99e',
'saab 99le', 'saab 99gle', 'subaru', 'subaru dl', 'subaru brz',
'subaru baja', 'subaru r1', 'subaru r2', 'subaru trezia',
'subaru tribeca', 'toyota corona mark ii', 'toyota corona',
'toyota corolla 1200', 'toyota corona hardtop',
'toyota corolla 1600 (sw)', 'toyota carina', 'toyota mark ii',
'toyota corolla', 'toyota corolla liftback',
'toyota celica gt liftback', 'toyota corolla tercel',
'toyota corona liftback', 'toyota starlet', 'toyota tercel',
'toyota cressida', 'toyota celica gt', 'toyouta tercel',
'vokswagen rabbit', 'volkswagen 1131 deluxe sedan',
'volkswagen model 111', 'volkswagen type 3', 'volkswagen 411 (sw)',
'volkswagen super beetle', 'volkswagen dasher', 'vw dasher',
'vw rabbit', 'volkswagen rabbit', 'volkswagen rabbit custom',
'volvo 145e (sw)', 'volvo 144ea', 'volvo 244dl', 'volvo 245',
'volvo 264gl', 'volvo diesel', 'volvo 246'], dtype=object)
```

```
[18]: df2=dataset.copy()
```

```
[19]: df2['company'] = df2['CarName'].str.split(" ", expand=True)[0]
df2['company'] = df2['company'].replace({'toyouta': 'Toyota', 'vw':
    ↳ 'Volkswagen', 'vokswagen': 'Volkswagen',
    'maxda':
    ↳ 'Mazda', 'porcshce': 'Porsche'})
df2['company'] = df2['company'].str.title()
df2['company'].value_counts()
```

```
[19]: Toyota      32
      Nissan      18
      Mazda      17
```


Mitsubishi	13
Honda	13
Volkswagen	12
Subaru	12
Peugeot	11
Volvo	11
Dodge	9
Buick	8
Bmw	8
Audi	7
Plymouth	7
Saab	6
Porsche	5
Isuzu	4
Jaguar	3
Chevrolet	3
Alfa-Romero	3
Renault	2
Mercury	1

Name: company, dtype: int64

```
[20]: plt.figure(figsize=(20, 6))

df_autox = pd.DataFrame(df2.groupby(['company'])['price'].mean().
    ↪sort_values(ascending = False))
df_autox.plot.bar()
plt.title('Car Company Name vs Average Price')
plt.show()
```

<Figure size 1440x432 with 0 Axes>



```
[21]: df_autox.rename(columns={'price': 'mean_price'}, inplace=True)
```

```
[22]: df_autox.head(10)
```

```
[22]:
```

company	mean_price
Jaguar	34600.000000
Buick	33647.000000
Porsche	31400.500000
Bmw	26118.750000
Volvo	18063.181818
Audi	17859.166714
Mercury	16503.000000
Alfa-Romero	15498.333333
Peugeot	15489.090909
Saab	15223.333333

```
[23]: dataset["engine_location"].value_counts()
```

```
[23]: front    202
      rear      3
      Name: enginelocation, dtype: int64
```

```
[24]: dataset.fuelsystem.value_counts()
```

```
[24]: mpfi      94
      2bbl      66
      idi       20
      1bbl      11
      spdi       9
      4bbl       3
      mfi        1
      spfi        1
      Name: fuelsystem, dtype: int64
```

## 10 Encoding categorical variable

```
[52]: dataset_new = dataset.copy()
```

```
[53]: encoders_nums = {"fueltype":{"diesel":1,"gas":0},
                        "aspiration":{"turbo":1,"std":0},
                        "doornumber": {"four": 4, "two": 2},
                        "drivewheel":{"fwd":0,"4wd":0,"rwd":1},
                        "cylindernumber":{"four": 4, "six": 6, "five": 5, "eight": 8,
                                           "two": 2, "twelve": 12, "three":3 },
                        "enginelocation": {"front":0,"back":1}
                        }

dataset_new = dataset_new.replace(encoders_nums)
```

```
[54]: dataset_new = pd.get_dummies(dataset_new, columns=["carbody",
↳ "engine", "fuelsystem", "enginelocation"], prefix=["body",
↳ "etype", "fsystem", "e_location"])
```

```
[65]: dataset_new = dataset_new.drop(["CarName"], axis=1)
```

```
[66]: dataset_new.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 43 columns):
#   Column                Non-Null Count  Dtype
---  -
0   car_ID                205 non-null   int64
```

1	symboling	205 non-null	int64
2	fueltype	205 non-null	int64
3	aspiration	205 non-null	int64
4	doornumber	205 non-null	int64
5	drivewheel	205 non-null	int64
6	wheelbase	205 non-null	float64
7	carlength	205 non-null	float64
8	carwidth	205 non-null	float64
9	carheight	205 non-null	float64
10	curbweight	205 non-null	int64
11	cylindernumber	205 non-null	int64
12	enginesize	205 non-null	int64
13	boreratio	205 non-null	float64
14	stroke	205 non-null	float64
15	compressionratio	205 non-null	float64
16	horsepower	205 non-null	int64
17	peakrpm	205 non-null	int64
18	citympg	205 non-null	int64
19	highwaympg	205 non-null	int64
20	price	205 non-null	float64
21	body_convertible	205 non-null	uint8
22	body_hardtop	205 non-null	uint8
23	body_hatchback	205 non-null	uint8
24	body_sedan	205 non-null	uint8
25	body_wagon	205 non-null	uint8
26	etype_dohc	205 non-null	uint8
27	etype_dohcv	205 non-null	uint8
28	etype_l	205 non-null	uint8
29	etype_ohc	205 non-null	uint8
30	etype_ohcf	205 non-null	uint8
31	etype_ohcv	205 non-null	uint8
32	etype_rotor	205 non-null	uint8
33	fsystem_1bbl	205 non-null	uint8
34	fsystem_2bbl	205 non-null	uint8
35	fsystem_4bbl	205 non-null	uint8
36	fsystem_idi	205 non-null	uint8
37	fsystem_mfi	205 non-null	uint8
38	fsystem_mphi	205 non-null	uint8
39	fsystem_spdi	205 non-null	uint8
40	fsystem_spfi	205 non-null	uint8
41	e_location_0	205 non-null	uint8
42	e_location_rear	205 non-null	uint8

dtypes: float64(8), int64(13), uint8(22)

memory usage: 38.2 KB

```
[67]: X=dataset_new.drop(["price"],axis=1)
      y=dataset_new["price"]
```



```
[68]: len(dataset_new.columns)
```

```
[68]: 43
```

## 11 Step-6: Dataset Splitting

```
[69]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split( X,y , test_size = 0.2,
↳random_state = 0)
print(X_train.shape)
print(X_test.shape)
```

```
(164, 42)
```

```
(41, 42)
```

## 12 Step-7: Model Implementation

- (1) Apply Linear regression,Lasso,Ridge and elastic net regression on the model.
- (2) Evaluate performance of different model.
- (3) Comment your observation

```
[70]: from sklearn.linear_model import LinearRegression

reg = LinearRegression().fit(X_train, y_train)
```

```
[72]: reg.score(X_train, y_train)
```

```
[72]: 0.956070774146671
```

```
[73]: reg.coef_
```

```
[73]: array([-1.39834906e+01,  5.65990898e+02,  6.15090102e+03,  1.89541668e+03,
         5.55971157e+02,  2.53161175e+03,  2.36765487e+02, -4.09007253e+01,
         5.63646392e+02,  1.84383691e+02,  5.07308735e+00,  1.33461886e+03,
         7.26129191e+01, -1.82848445e+03, -2.44930615e+03, -9.56663239e+02,
        -1.86609539e+01,  1.61651527e+00,  5.25998459e+01,  8.13788134e+01,
         3.41415840e+03, -7.23092654e+02, -4.56745994e+02, -3.04057644e+02,
        -1.93026211e+03, -7.37875486e+02,  0.00000000e+00, -6.18974485e+03,
         5.20783434e+02,  1.00372714e+03, -3.55651366e+03,  8.95962342e+03,
        -3.56444407e+01, -8.12430759e+00, -1.53848439e+03,  6.15090102e+03,
        -2.10733609e+03,  7.47346167e+02, -1.23721227e+03, -1.97144569e+03,
        -7.18110754e+03,  7.18110754e+03])
```

```
[74]: len(reg.coef_)
```

```
[74]: 42
```

```
[75]: reg.intercept_
```

```
[75]: -60769.327782530134
```

```
[76]: y_pred = reg.predict(X_test)
```

```
[78]: from sklearn.metrics import mean_squared_error
```

```
MSE = mean_squared_error(y_test,y_pred)
```

```
print("MSE :" , MSE)
```

```
RMSE = np.sqrt(MSE)
```

```
print("RMSE :" ,RMSE)
```

```
MSE : 14490755.960495027
```

```
RMSE : 3806.6725575619225
```

```
[79]: from sklearn.metrics import r2_score
```

```
r2 = r2_score(y_test,y_pred)
```

```
print("R2 :" ,r2)
```

```
R2 : 0.8128204904846368
```

```
[80]: from sklearn.linear_model import Lasso
```

```
lasso = Lasso(alpha=0.0001 , max_iter= 3000)
```

```
lasso.fit(X_train, y_train)
```

```
[80]: Lasso(alpha=0.0001, max_iter=3000)
```

```
[81]: lasso.score(X_train, y_train)
```

```
[81]: 0.9560707741460119
```

```
[82]: lasso.coef_
```

```
[82]: array([-1.39835200e+01,  5.65989390e+02,  3.32534801e+03,  1.89541830e+03,  
        5.55970164e+02,  2.53160999e+03,  2.36765298e+02, -4.09011240e+01,  
        5.63647213e+02,  1.84384363e+02,  5.07308389e+00,  1.33461104e+03,  
        7.26130345e+01, -1.82848763e+03, -2.44929990e+03, -9.56649538e+02,  
       -1.86608460e+01,  1.61651101e+00,  5.25992111e+01,  8.13787277e+01,  
        4.25450742e+03,  1.17255303e+02,  3.83606606e+02,  5.36295892e+02,  
       -1.08990754e+03, -8.94070556e+02,  0.00000000e+00, -6.34592760e+03,
```

```

3.64586780e+02, 8.47534733e+02, -3.71270453e+03, 8.80337191e+03,
-1.44916536e+03, -1.42164071e+03, -2.95195558e+03, 7.56275630e+03,
-3.52082787e+03, -6.66167590e+02, -2.65071430e+03, -3.38494725e+03,
-1.43621992e+04, 9.60023061e-09])

```

```

[83]: from sklearn.model_selection import GridSearchCV
lasso = Lasso()
parameters = {'alpha': [
    ↪[1e-15,1e-13,1e-10,1e-8,1e-5,1e-4,1e-3,1e-2,1e-1,1,5,10,20,30,40,45,50,55,60,100,0.
    ↪0014]}
lasso_regressor = GridSearchCV(lasso, parameters,
    ↪scoring='neg_mean_squared_error', cv=5)
lasso_regressor.fit(X_train, y_train)

```

```

[83]: GridSearchCV(cv=5, estimator=Lasso(),
    param_grid={'alpha': [1e-15, 1e-13, 1e-10, 1e-08, 1e-05, 0.0001,
    0.001, 0.01, 0.1, 1, 5, 10, 20, 30, 40, 45,
    50, 55, 60, 100, 0.0014]}},
    scoring='neg_mean_squared_error')

```

```

[84]: print("The best fit alpha value is found out to be :",lasso_regressor.
    ↪best_params_)
print("\nUsing ",lasso_regressor.best_params_, " the negative mean squared
    ↪error is: ", lasso_regressor.best_score_)

```

The best fit alpha value is found out to be : {'alpha': 1}

Using {'alpha': 1} the negative mean squared error is: -5715234.081444084

```

[86]: y_pred_lasso = lasso_regressor.predict(X_test)

```

```

[87]: MSE = mean_squared_error(y_test,y_pred_lasso)
print("MSE :", MSE)

RMSE = np.sqrt(MSE)
print("RMSE :",RMSE)

r2 = r2_score(y_test,y_pred_lasso)
print("R2 :",r2)

```

MSE : 14370473.87325172  
 RMSE : 3790.8407871146105  
 R2 : 0.8143741942496484

```

[88]: from sklearn.linear_model import Ridge
ridge = Ridge()

```

```

parameters = {'alpha':  

    ↳[1e-15,1e-10,1e-8,1e-5,1e-4,1e-3,1e-2,1,5,10,20,30,40,45,50,55,60,100]}
ridge_regressor = GridSearchCV(ridge, parameters,  

    ↳scoring='neg_mean_squared_error', cv=3)
ridge_regressor.fit(X_train,y_train)

```

```

[88]: GridSearchCV(cv=3, estimator=Ridge(),
        param_grid={'alpha': [1e-15, 1e-10, 1e-08, 1e-05, 0.0001, 0.001,
                                0.01, 1, 5, 10, 20, 30, 40, 45, 50, 55, 60,
                                100]},
        scoring='neg_mean_squared_error')

```

```

[89]: print("The best fit alpha value is found out to be :",ridge_regressor.
        ↳best_params_)
print("\nUsing ",ridge_regressor.best_params_, " the negative mean squared_
        ↳error is: ", ridge_regressor.best_score_)

```

The best fit alpha value is found out to be : {'alpha': 1e-08}

Using {'alpha': 1e-08} the negative mean squared error is: -6297757.865667018

```

[90]: y_pred_ridge = ridge_regressor.predict(X_test)

```

```

[92]: MSE = mean_squared_error(y_test,y_pred_ridge)
print("MSE :", MSE)

RMSE = np.sqrt(MSE)
print("RMSE :",RMSE)

r2 = r2_score(y_test,y_pred_ridge)
print("R2 :",r2)

```

MSE : 14490753.857985102  
 RMSE : 3806.672281400791  
 R2 : 0.8128205176431078

```

[93]: from sklearn.linear_model import ElasticNet
        #a * L1 + b * L2
        #alpha = a + b and l1_ratio = a / (a + b)
        elasticnet = ElasticNet(alpha=0.1, l1_ratio=0.5)

```

```

[94]: elasticnet.fit(X_train,y_train)

```

```

[94]: ElasticNet(alpha=0.1)

```

```

[95]: elasticnet.score(X_train, y_train)

```

```
[95]: 0.9215222910941743
```

```
[96]: y_pred_en = elasticnet.predict(X_test)
```

```
[98]: MSE = mean_squared_error(y_test,y_pred_en)
print("MSE :", MSE)
RMSE = np.sqrt(MSE)
print("RMSE :", RMSE)
r2 = r2_score(y_test,y_pred_en)
print("R2 :", r2)
```

```
MSE : 12008339.137873283
```

```
RMSE : 3465.305056971649
```

```
R2 : 0.8448862822582186
```